A Case Study in Preserving a High Energy Physics Application

Haiyan Meng, Matthias Wolf, Peter Ivie, Anna Woodard, Michael Hildreth, and Douglas Thain
Department of Physics and Department of Computer Science and Engineering
University of Notre Dame
{hmeng|mwolf3|pivie|awoodard|mhildreth|dthain}@nd.edu

ABSTRACT

The reproducibility of scientific results increasingly depends upon the preservation of computational artifacts. Although preserving a computation to be used later sounds easy, it is surprisingly difficult due to the complexity of existing software and systems. Implicit dependencies, networked resources, and shifting compatibility all conspire to break applications that appear to work well. To investigate these issues, we present a case study of a complex high energy physics application. We analyze the application and attempt several methods at extracting its dependencies for the purposes of preservation. We report on the completeness, performance, and efficiency of each technique, and offer some guidance for future work in application preservation.

1. INTRODUCTION

Reproducibility is a cornerstone of the scientific process. In order to understand, verify, and build upon previous work, one must be able to first recreate previous results by applying the same methods. Historically, reproducibility this has been accomplished through painstaking detailed documentation recorded in lab notebooks, which are then summarized in peer-reviewed publications. But as science increasingly depends on computation, reproducibility must also encompass the environments, data, and software involved in each result. It is widely recognized that informal descriptions of software and systems – although common – are insufficient for reproducing a computational result accurately. A more automated and comprehensive approach is required.

The overall reproduction of a computation has three broad components, each of which suggests somewhat different approaches:

- The **computing environment**, consisting of the basic hardware and the operating system can be preserved as physical artifacts or as a combination of virtual machine monitor (hardware) and virtual machine image (operating system).

- The **scientific data** to be analyzed has historically received the most attention for curation. In a large, well-organized project, it may be stored in a data repository or database management system, with associated documentation and a curation strategy. In a small effort, it could simply be a handful of files.

- The **software environment** includes the source code, binaries, scripts, configuration files, and everything else needed to execute the desired code. As with data, the software could be drawn from a well-managed software repository, or it could be a handful custom scripts that exist in the user's home directory.

In a very abstract sense, reproducing a computation is trivial. Assuming a computation is deterministic, one must simply preserve all of the inputs to a computation, then re-run the same code in an equivalent environment, and the same result will be produced. For a small custom application on a modest amount of data, this could be accomplished by capturing the environment, data, and software within a single virtual machine image, and then depositing the virtual it into a curated environment. The publication could then simply refer to the identifier of the image, which the interested reader can obtain and re-use. This approach has been used to some success with systems X, Y, and Z. [1]

However, this simple approach is not sufficient for large applications that are run in complex social environments.

- There may be **implicit dependencies** on items that are not apparent to the end user. For example, they may understand that they rely on a particular data analysis package, but would have no reason to know that the package has further dependencies on other libraries and configuration files. Or, they may know that the computation only runs correctly on a particular machine, but not know this is because it relies on data in a filesystem that is mounted only on that machine.

---

[1]Of course, we are glossing over the problem that hardware architectures and virtual machines also change, so one must also preserve the VMM software necessary to run the image. The VMM itself depends on a software environment which must also be preserved. A long-term preservation system might end up running a whole stack of nested virtual machines in order to provide the desired environment!
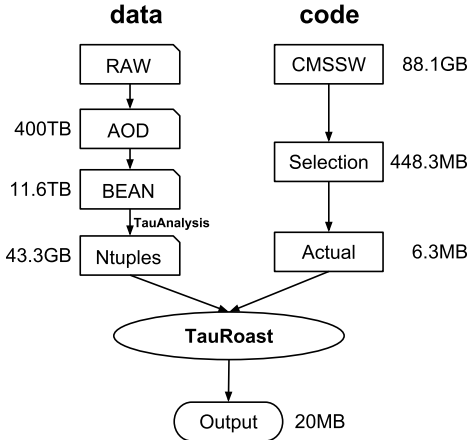
data   **code**

RAW

400TB   AOD   CMSSW   88.1GB

11.6TB   BEAN   Selection   448.3MB

        TauAnalysis

43.3GB   Ntuples   Actual   6.3MB

            TauRoast

            Output   20MB

Figure 1: Inputs to Tau Roast

Name	Location	Total	Named	Used
Ntuples data	hdfs	24TB	43.3GB	20GB
CMSSW binaries	local	88.1GB	448.3MB	6.3MB
Tau source	git	73.7MB	73.7MB	73.7MB
PyYAML binaries	http	52MB	51MB	51MB
.h file	http	41KB	41KB	41KB
Misc commands	PanFS	155TB	N/A	1.6MB
Linux commands	localFS	110GB	N/A	68.3MB
Cern files	CVMFS	7.4GB	N/A	103MB
NDCMS	NFS	862GB	N/A	53KB
Local store	AFS	10.2GB	N/A	34MB
Total		180.1TB	N/A	21GB

The first column illustrates the total size of each data and software source; the second column illustrates the size of the named files from each source; the third column illustrates the size of acutally used data from each source. N/A denotes it is hard to figure out the named size of implicit dependencies directly.

Table 1: Data and Code Used by Tau Roast

- The **granularity** of the dependencies may not be well understood. For example, the user may understand that a computation depends upon a data collection that is 1TB in overall size, but not have detailed knowledge that it only requires three files totalling 300MB out of that whole collection

- There may be dependencies upon **networked resources** that are inherently external to the system, such as a database, a code repository, or a scalable filesystem. For such resources, it must be decided whether the dependency will simply be noted, or if it must be incorporated whole or in part.

- Where **common dependencies** are widely used, it may be ineffecient or impossible to store one copy of each dependency for each archived object. Some form of sharing or de-duplication is necessary in order to keep the archive to a reasonable size.

We do not claim to have solved these problems in any comprehensive way. Rather, our aim in this paper is to highlight the scope of the problems by presenting a case study of one complex application. The application is presented to us first in the form of an email that describes in prose how to install the software and run the analysis. We perform several successive refinements to convert it into an executable and preservable object. Then, we develop two techniques for observing and capturing the dependencies associated with the system, comparing the cost of capture, the size of the preserved object, and the flexibility of the resulting object. We describe how each of these techniques may interact with a future archive of preserved software artifacts, and conclude with some reflections on the challenges of preservation and advice for future efforts.

2. OVERVIEW OF TAU ANALYSIS

Within the ongoing investigation of the Higgs boson at the CMS detector, part of the LHC at CERN, the Higgs production in association with two top quarks allows to measure the Higgs coupling strength to top quarks. As the Higgs boson is to short-lived to be detected itself, it has to be reconstructed from its decay products.

The application which is the study of this paper is called *TauRoast* It searches for cases where the the Higgs boson decays to two tau leptons. The leptons are not observed directly, but by the particle showers that they generate. So, the analysis must search for detector events that show a signature of decay products compatible with both hadronic tau and top decays. Properties of such events are used to distinguish the events of interest (Higgs decays) from all other events and are also used in further statistical analysis.

Figure 1 shows that both the code and data that form *TauRoast* are drawn from large repositories through multiple steps of reduction.

Data Sources. The CMS collaboration provides analysis end-users with a pre-processed and reduced data format, AOD, containing information for events, i.e., proton-proton collisions with a signature of interest, in the form of reconstructed particles. This format is based on the RAW output of the CMS detector readout electronics and reconstructed world-wide. Both real and simulated data are available for examination.

As AOD data are too large to be iteratively processed repetitively in an physics analysis workflow, it is normally reduced further in structural complexity and content. For the analysis under investigation here, this is a two-step process. First, the AOD data are processed at the Notre Dame working group cluster to BEAN events, containing only trivial data containers packed in vectors. This step is time and CPU intensive and its output contains data of 11.6 TB to be analyzed by the tau analysis. It is performed by a small custom code framework, which is built on top of the CMS software stack, CMSSW, and uses packages provided by several other special interest groups within CMS. While the CMSSW framework is installed locally, the various packages used are checked out from CVS, and the BEAN framework is stored in git. This is scheduled to change, as the CMSSW distribution model switches to a virtual filesystem mounted via FUSE, and special interest groups move their code to git. The BEAN format, production code, and data are shared

within the analysis group looking at Higgs production in association with top quarks, which is formed by groups from a few American and European universities, consisting of up to a few dozen contributors.

In the second step, which is the beginning of the actual tau analysis, the data are reduced to variables relevant to the tau roast procedure, while only events matching basic quality criteria are kept. This results in a dataset of 43.3 GB. Again, the Notre Dame CMS groups cluster resources are used to perform this reduction and selection, running highly customized software, built on CMSSW and the BEAN framework, but with output code written and maintained by a few people only. Again, the code is stored in a git repository.

The final data analysis, investigated below, can be run as a single process, and contains a stringent event selection to keep only high quality candidate events for the underlying physical process (using about 20 MiB of space). Quantities from the selected events can be both plotted and used in multivariate analysis to determine the level of expected signal in real data. This package is written using the CMSSW build framework, but only utilizes code from ROOT, a particle physics toolkit underlying CMSSW, and a few external python dependencies for convenience. The latter have to be manually fetched and installed, while the analysis program is built by CMSSW after being checked out of git.

Code Sources. Like many scientific codes, the central algorithm of *TauRoast* is expressed in a relatively small amount of custom code developed by the primary author. But, the code cannot run at all without making use of an enormous collection of software dependencies. Some of these dependencies are standard to operating systems worldwide, some are standardized across the entire high-energy physics field, some are particular to small collaborative groups, and a few are very specific to a single researcher.

The largest of these repositories is the CMS Software Distribution (CMSSW), a carefully-curated selection of software packages which is distributed in binary form. Historically, CMSSW was downloaded and installed on shared filesystems within HPC centers. In recent years, distribution has moved to an on-demand delivery system known as CVMFS [3]. The content of CMSSW is managed very carefully by a centralized team whose main goal is to ensure that the current version of the software operates correctly on the operating systems and architectures currently in use. However, preservation is not an objective of the group, and so there is no guarantee that old versions of CMSSW operate in new environments, or vice versa.

TauRoast was provided to us in the form of an email which described, in prose, how to obtain the source, build the program, and run it correctly on one specific machine at Notre Dame, with no particularly guarantee that it will run anywhere else in the world. Although this starting point may seem extreme, it is perfectly natural for collaborators to share configurations with each other in this form, and to rely on the presence of a working environment with appropriate dependencies already installed.

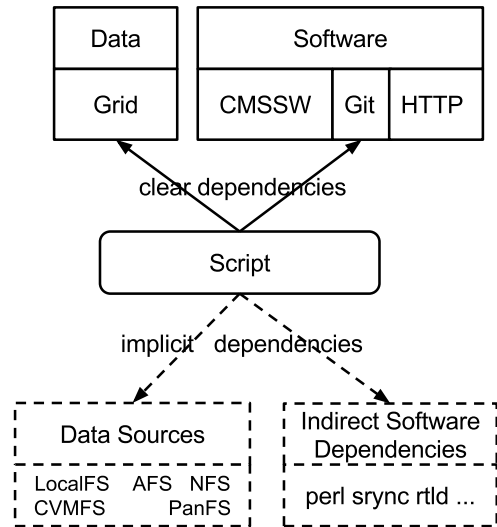The authors played the role of curators, whose job it is to

Figure 2: Dependencies

prepare the application for permanent archival by determining and packaging the dependencies.

2.1 Workflow of Matthias's Example

The workflow of Matthias's exmaple can be divided into eight procedures.

(a) Declare environment variables

(b) Obtain software from CMSSW

(c) Obtain software from Git

(d) Obtain software from some public HTTP web links

(e) Obtain software from one private home page

(f) Build software enviroment using SCRAM and Python

(g) Install grid control, obtain CMS data from grid and store the data into HDFS mounted as one local file system.

(h) Actual data analysis.

2.2 Key Observations of the Example

2.2.1 Complexity of Data and Software Dependencies

Figure 2 illustrates all the dependencies involved in the program, which can be divided into two categories: clear dependencies and implicit dependencies. Clear dependencies include scientific data dependencies and direct software dependencies which can be extracted directly from the analysis program. The scientific data of this example comes from the Grid. Direct software dependencies include CMSSW, Git, HTTP and one personal home page.

Implicit dependencies include indirect software dependencies which is necessary for the successful execution of direct software dependencies and other implicit dependencies.

For example, we know Git is one direct software dependency from the analysis program, which further depends on perl, python, rsync, openssh-clients and other packages. The analysis program accesses one file called File1, which is one symbolic link file further referring to another file named File2. The implicit dependencies also include some remote file systems mounted as local file systems, such CVMFS, AFS, NFS, PanFS and HDFS. More importantly, the successful execution of this program greatly depends on the underlying OS and hardware platform.

2.2.2 Size of Data and Software Dependencies

The size of data and software dependencies is astonishing. The size of BEAN is 11.6TB, which is too large to be analyzed directly by the tau roast program. After one preprocessing which only collects the events matching basic quality criteria, the data size is still very large, 43.3GB.

The first column of Table 1 illustrates the size of each source of data and software. The size of PanFS is as large as 155TB and the size of Ntuples data is 24TB. The total size of all the dependencies is about 181TB, which directly makes the idea of preserving the original execution environment completely impossible.

2.2.3 Complexity of the Program

Complex data and software dependencies, complex environment building process, the large size of data and software involved in the tau roast program, together with the access authority problem, determine the complexity of the program. In order to obtain the experimental data from the Grid, the author needs to install and config grid-control software. The access authority of the Grid is also necessary for the data acquisition. The software acquisition from CMSSW, Git and HTTP and environment building process is time-consuming. The execution time of actual data analysis is about 20 minutes. However, the software acquisition and environment building time, excluding data acquisition time, is about 14 minutes. If the acquisition of experimental data is considered, actual data analysis only occupies one small percentage of the whole time consumption. The size of data and software dependencies further complicates the program.

2.2.4 Really Used Data and Software Size

Even if the orginal data size referred by the program is astonishing, the size of really used data and software is greatly smaller. Figure 1 illustrates the inputs of the tau roast program. After the TauAnalysis procedure, the data size of Ntuples is 43.3GB, however, the atual size of really used data in the program is 20GB. The same trend occurs to the size of software and other dependencies. The size of CMSSW repository is 13.1GB, the size of packages checked out from CMSSW is 448.3MB, and the size of actually used CMSSW files is only 6.3MB. The total size of PanFS mounted as `/pscratch` is 155TB, however, the really used data from `/pscratch` is only 1.6MB. The third column of Table 1 illustrates the size of actually used data and software from each dependency.

2.2.5 Stability of Dependencies

The program involves obtaining software from the home page of another graduate student from University of Notre Dame. However, the maintenance of this home page will terminate after the graduation of the student. The stability of each dependency must be evaluated to ensure the program can be repeated by others.

The stability of the Grid is relatively higher. However, the public web resources and personal home page is not stable. Too many factors like the termination of one web site, the graduation of the home page's author, would result in the failure of resource access. CMSSW, which seems to be stable, in fact, also involves stability problems. CMSSW provides different collections of accessible versions to different architectures. Transformation from one architecture to another may result in the failure of accessing one certain CMSSW version. At the beginning of our reseach, we can use the cvs command to check out packages from CMSSW anonymously. However, this anonymous access has been cancelled now.

3. PRESERVATION STRATEGIES

Figure 3 illustrates the evolution history of different preservation strategies. First, an email depicting how to repeat the experiment is used. Then, all the necessary steps to reproduce the experiment mentioned in the mail are collected into one shell script. In order to decouple data sources and data references, one formulated script is introduced to express all the dependencies at the beginning of the script in the form of environment variables. Finally, this decoupling idea is further extended into one map file which keeps the relationship between the data reference in the script and the actual storage location. One independent package containing all the data and software dependencies is generated to make the reproduction process more easier.

3.1 Solution 1: Email

In order to repeat Matthais's example, we consulted the original author by email about the necessary work for the experiment. In response, Matthias introduced the general workflow of his tau roast program through one long email including notes, linux shell commands, web linksi, as illustrated by the Version 1 of Table 2.

However, this solution to repeat one experiment has three potential drawbacks. Firstly, the experience of repeating one tau roast program through emails is chaotic. You need to constantly jump around multiple web links. There are overlap between the content of the email and the content of web links, which needs the new user to merge them. Multiple communication through emails is necessary to ensure the successful reproduction of the whole analysis. For example, the original email refers to one environment variable called CMSSW_base without clear declaration, the new user needs to send one email to the original writer to obtain its accurate value. Secondly, the necessary procedures to repeat the experiment, including software acquisition from different sources, is complex for the new user. In Matthias's example, the sources of software includes CMSSW, Git, HTTP. The access of CMSSW requires the new user to be an authorized user of CMSSW. What's more, some parts of the workflow are unrepeatable. The third step of this experiment requires the new user to own the access authority of the Grid and involves the analysis of AOD data with the size of about 400TB.
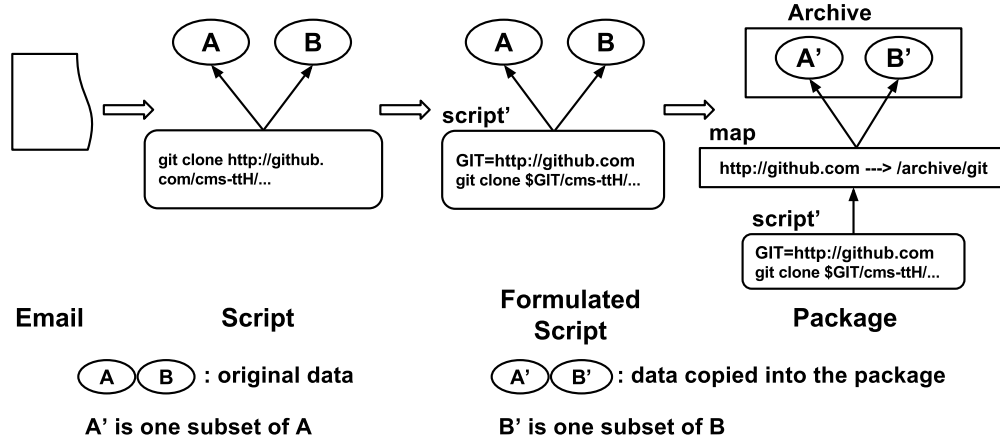
Figure 3: Version Evolution

Implication: Directly repeating the experiment using the workflow description and results provided by the original author is complex and difficult. Some extra work must be done to make the reproduction process easier.

3.2 Solution 2: Script

One possible solution for the access authority problem of grid data is to grant the new user the authority to access gird data directly. Another possible solution is let the new user directly operate on the machine the original author used. As for the complexity of jumping between multiple web links, people may suggest that the original author should generate one script including all the contents from different web links.

To test out these possible solutions, we integrate the content of all the notes, commands and web links involved in the emails into one neat, complete shell script, which begins with the definition of environment variables, software acquisition from CMSSW, Git and other web resources, and software installation, ends with the execution of the actual analysis program. Version 2 of Table 2 illstrates the merged script. As for the data acquisition from the Grid, we directly use the local copy in HDFS to avoid requiring the new user to obtain a grid certification.

Because the script includes all the necessary procedures for the reproduction of one analysis program, the readability and friendliness of this solution is higher than that of the email format.

3.3 Solution 3: Formulated Script

The script can reduce the complexity of repeating one experiment through the integration of all the necessary procedures. However, the data and software dependencies are still randomly distributed across the whole script, which requires one complete scan of the script to obtain the dependency list.

To decouple data sources and data references and make the data dependencies clear, Version 2 of Table 2 is formulated into one new version, as shown in the Version 3 of Table 2. Each new environment variable at the beginning part of the script is corresponding to one dependency. All the follow-

ing access to data or software dependency will refer to its corresponding environment variable. All the environment variables of dependencies form one map, which maintains the target address of each dependency.

This script style may function as one new guideline to the original author of one experiment, which expresses the dependencies more clearly and expedites the preparation process to repeat one experiment. The introduction of the map file also reduces the workload of changing one dependency. For example, if we want to utilize one new git package to analyze the same dataset, only the modification of environment variable corresponding to the original git package is necessary. Without the map mechanism, the whole script needs to be scanned to figure out and replace all the references of the original git package.

So far, the data access authority is ignored and the reproduction of the experiment is on the machine the original author used. As for grid data, using the data copy stored in the machine where Matthias executed the experiment requires the new user to have the authority to access the machine, which complicates the management of the original machine, even is impossible if the original machine executes rigid user access control. Granting everyone who wants to repeat the analysis the access authority for grid is unacceptable to system administrator.

Implication: All the data and software involved in the experiment should be provided to the new user in the format of one self-contained package so that the new user can avoid access authority problems. In addition, the requirements of the underlying OS and hardware should also be figured out and provided to the new user.

3.4 Solution 4: Fine-Grained Toolkit - Package

The difficulty of data access authority acquisition enforces us to find out one solution, in which the reproduction of the original analysis can be done without any external dependency. That is, one independent and self-contained package containing all the data and software dependencies is neces-

Version 1: Email

1. Create a CMS release,
     e.g. `cmsrel CMSSW_5_3_11_patch3`
 2. Install the BEAN packages as the instructions:
     `https://github.com/cms-ttH/BEAN/blob/...`
 3. Install grid-control:
     `svn co https://ekptrac.physik.uni-ka/...`
 4. INstall the TauAnalysis package:
     `git clone https://github.com/matz-e/...`
     `scram b -j32`
 5. Fix grid_control.cfg and run it.
 6. Perform the actual tau roast program.

Version 2: Script

`set CMSSW_BASE = (CMSSW_5_3_11_patch3)`
 `cmsrel $HOME/$CMSSW_BASE`
 `cvs co -r V03-09-23 PhysicsTools/PatUtils`
 `git clone https://github.com/cms-ttH/BEAN.git`
 `wget -r http://nd.edu/~abrinke1/...`
 `scram b -j32`
 `wget http://pyyaml.org/download/pyyaml/PyYAML...`
 `#the experiment data is from HDFS`
 `cd $HOME/$CMSSW_BASE/src/PyYAML-3.10`
 `cmsenv`
 `python setup.py install -user`
 `scripts/roaster data/generic_ttl.yaml`

version 3: Formulated Script

`set CMSSW_BASE = (CMSSW_5_3_11_patch3)`
 `set ` **`GIT`** ` = (https://github.com)`
 `set ` **`PYYAML`** ` = (http://pyyaml.org)`
 `set ` **`ND`** ` = (http://nd.edu)`
 `cmsrel $HOME/$CMSSW_BASE`
 `cvs co -r V03-09-23 PhysicsTools/PatUtils`
 `git clone $`**`GIT`**`/cms-ttH/BEAN.git`
 `wget -r $`**`ND`**`/~abrinke1/ElectronEffectiveArea.h`
 `scram b -j32`
 `wget $`**`PYYAML`**`/download/pyyaml/PyYAML...`
 `#the experiment data is from HDFS`
 `cd $HOME/$CMSSW_BASE/src/PyYAML-3.10`
 `cmsenv`
 `python setup.py install -user`
 `scripts/roaster data/generic_ttl.yaml`

Version 4: Fine-Grained Toolkit - Package

`set CMSSW_BASE = (CMSSW_5_3_11_patch3)`
 `set ` **`GIT`** ` = (https://github.com)`
 `set ` **`PYYAML`** ` = (http://pyyaml.org)`
 `set ` **`ND`** ` = (http://nd.edu)`
 `cmsrel $HOME/$CMSSW_BASE`
 `cvs co -r V03-09-23 PhysicsTools/PatUtils`
 `git clone $`**`GIT`**`/cms-ttH/BEAN.git`
 `wget -r $`**`ND`**`/~abrinke1/ElectronEffectiveArea.h`
 `scram b -j32`
 `wget $`**`PYYAML`**`/download/pyyaml/PyYAML...`
 `#the experiment data is from HDFS`
 `cd $HOME/$CMSSW_BASE/src/PyYAML-3.10`
 `cmsenv`
 `python setup.py install -user`
 `scripts/roaster data/generic_ttl.yaml`

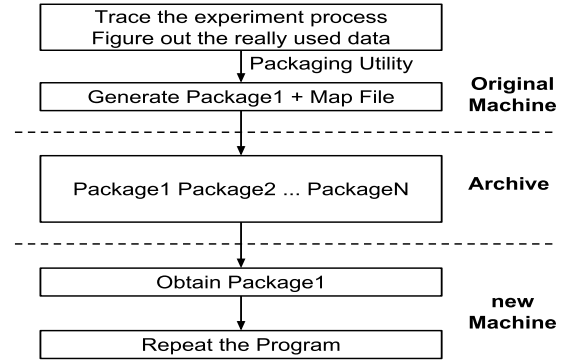**Table 2: Scripts of each Solution**

Figure 4: Relationship of Roles

sary.

Someone may suggest that it should be the responsibility of the original author to generate the required package. However, letting the original author provide the package. One reason is that figuring out the underlying dependencies of each software is complex and time-consuming and even impossible for the original author. In this experiment, the machine used for the experiment is one public machine of physics department, and Matthias is one common user without root authority. The underlying OS and supporting softwares are installed and maintained by the IT department of the university. On the other hand, the architecture design of the required package including all the data and software dependencies is not under the research field of physists.

To generate one self-contained and independent package for one experiment, all the clear and implicit dependencies must be figured out firstly. The map mechanism of Solution 3 can provide clear data and software dependencies. As for the implicit dependencies, directly preserving the whole OS together with the data stored on the disk of the original machine or the data from other filesystems mounted as local filesystems, such as HDFS and PanFS, is not feasible. One efficient mechanism which can figure out the really used parts of all these filesystems is necessary. The requirements of the underlying OS and hardware architecture can be easily found with system tools such as uname or lsb_release.

Then one new package including all the data and software dependencies can be generated and published. One packaging utility is necessary for the original author to generate the package. To multiplex the common data which is involved in different experiments, such as OS images and common software libraries, the data of different packages will be rearranged and categorized in the archive. One description file, which includes the experiment aim, dependency list, package size and relevant information, will be generated for each package.

The relationship of different roles involved in the experiment preservation and reproduction is shown in Figure 4. The original author is responsible for generating one package with the help of the packaging utility. Then the package, together with its map file and description file will be uploaded

Archive Path	Service Description
/archive/OS/path	OS Image
/archive/software/path	Common Software Library
/archive/CMSSW/path	CMS Software Library
/archive/git/path	Git Software Library
/archive/http/path	http resources
/archive/experiment/path	experiment private files
/archive/grid-data/path	data from the Grid

Table 3: Structure Organization of the Archive

into the archive. When another scholar wants to repeat the experiment, one copy of the package and its corresponding map file will be downloaded into the new machine.

Table 3 illustrates the structure of the remote archive. Item `/archive/OS` includes the images of different Operating Systems. Item `/archive/software` integrates all the commonly used software, such as git, python, perl. The archive also organizes the data from CMSSW, Git, HTTP and the Grid. Item `/archive/experiment` includes the all the experiments submitted into the archive. The private files and description files of each experiment will be organized together.

The Version 4 of Table 2 illustrates the script used in Solution 4, which is the same as the one used in Solution 3. However, one map file is necessary for the relocation of the data access targets, as show in Figure 3. The map file redirects the git access path into `/archive/git` from the original path (`http://github.com`) referred in the script. This design decouples the experiment script and the actual data access targets, which minimizes the impact of the evolution of different data dependencies and ensures the transparent access. The modification of the archive only introduces the minimal changes of the map file on the client side.

The archive supports two different data preservation models: Internal and External. Internal method will preserve the data in the archive. External method refuses to preserve the content of data, but only preserve the reference to the actual storage place of data. For example, the size of experiment data from the Grid is extremely large, and storing the same data in the archive is time-consuming and space-consuming. Through External method, the archive only preserves one reference to the data inside the remote Grid.

This solution tries to create one self-contained and independent package for each experiment and integrate different packages for different experiments into one archive to multiplex common data. One packaging utility, which can figure out all the really used data of one program will be provided to help the original author to generate the package. The archive can maintain the experiment dependencies through Internal or External method. The archive itself will be responsible for the data maintenance and relevant authority access problem. The new user can repeat one experiment through the interaction with the archive.

4. ONE IMPLEMENTION OF FINE-GRAINED TOOLKIT USING PARROT
The section introduces our implementation of one fine-grained packaging utility based on Solution 4. With the help of one

```
set CMSSW_BASE = (CMSSW_5_3_11_patch3)
cd $HOME/$CMSSW_BASE/src/PyYAML-3.10
cmsenv
python setup.py install -user
cd $HOME/$CMSSW_BASE/src/ttH/TauRoast
scripts/roaster data/generic_ttl.yaml
```

Table 4: Script of Fine-Grained Toolkit using Parrot

virtual filesystem access tool, Parrot, we trap all system calls of one program through ptrace debugging interface, and collect all the names of accessed files. Then one self-contained package containing all the accessed files is generated to help others repeat the experiment.

4.1 Working principle of Packaging Utility
Parrot is a virtual filesystem access tool which attaching existing programs to a variety of remote I/O systems including http, ftp, gridftp, irods, HDFS, xrootd, grow and chirp. It traps all system calls of one program through ptrace debugging interface, and replaces them with remote I/O operations as desired. Through executing one program under Parrot, all the paths of files involved in this program can be recorded.

With the help of Parrot, one packaging utility which generates one independent package for one program to make the reproduction of the program convenient can be deployed. The starting point of the packaging utility is one successful execution sandbox (the data from grid has been preserved in HDFS and the software from CMSSW, Git and HTTP has been on the local machine.). We re-execute the actual data analysis code under Parrot and get the name list of all the files actually accessed during the execution process of the actual data analysis. Then, according to the filename list, one package containing all the necessary data and software for one analysis program is generated. Next time, when another scholar wants to repeat the program, he only needs to obtain the package and directly execute the actual analysis program inside the package. Figure 4 illustrates the working principle of the packaging utility.

Table 4 illustrates the simplified script, which only contains the necessary environment variables and the actual analysis command.

4.2 Workflow of Packaging Utility
Figure 5 illustrates the workflow of this solution. Firstly, execute the analysis program under Parrot and obtain one filename list containing all the accessed files. Then generate one package containing all the files included in the filename list. Finally, rerun the program with the help of the package. Each step will be examined in details.

(1) execute one analysis program under Parrot and obtain filename list (L1) of it

Parrot is one virtual file system that can support user access to multiple underlying file systems. Parrot traps each system call involved in the process of data access, figures out the type of file system and redirects it into corresponding operations to the accessed file system. During this process,
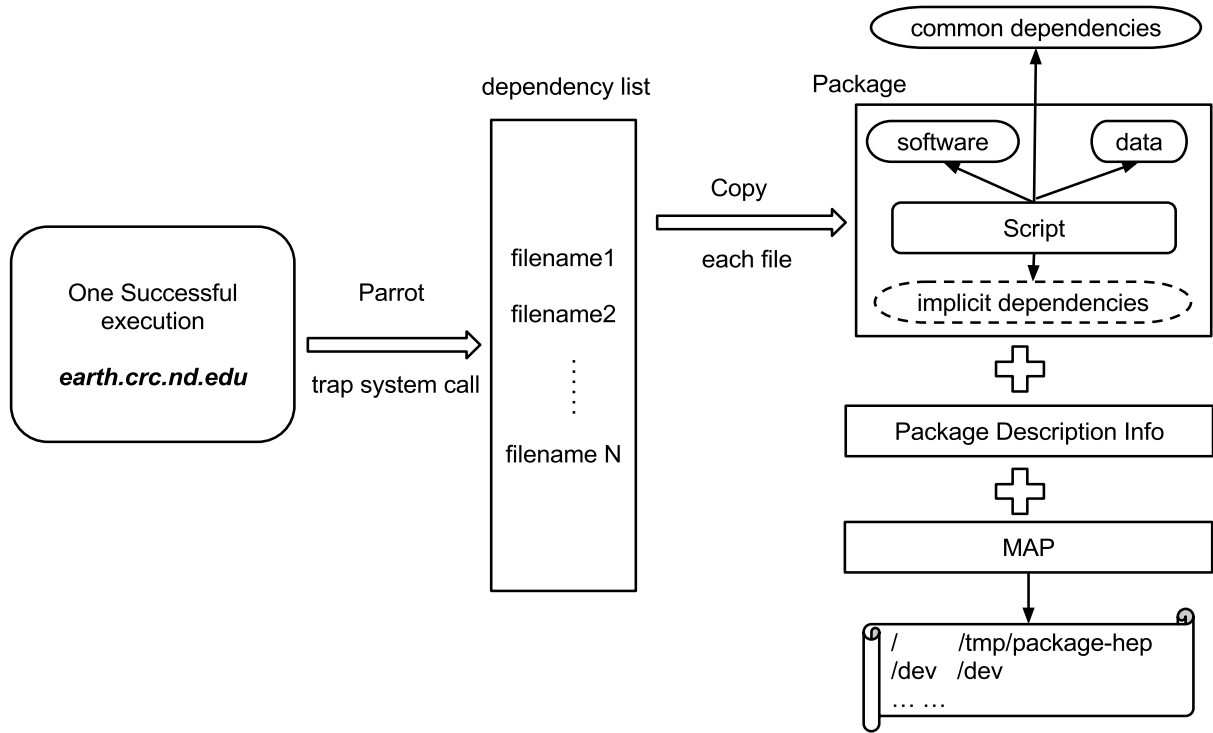
Figure 5: Workflow of Solution 4

each accessed file, together with the access type, such as open, stat, read and write, is recorded into one file.

The command used to generate the filename list is as follows. The -L parameter refers to the path of the file containing all the accessed file name of the experiment.

```
parrot_run -L namelist /bin/tcsh script_v4.sh
```

The original ouput of this command is one namelist file with the size of 6.6MB and 132,058 lines, each of which corresponds to one one accessed file of the program. We noticed that some items appeared multiple times in the filename list. To reduce the packaging time, de-duplication and sorting techoniques are used for the namelist file, generating one smaller file with the size of 3.2MB and 67,178 lines.

(2) generate one package containing the files of L1

The packaging utility iterates each filename inside L1 and copies it into the target package. Finally, the target package together with its description information, and one map file, which redirects the access of files from different file systems into the files inside the target package, will be generated as the output of this step. The packaging process also runs within Parrot environment to access data from different file systems.

The command used to generate the package is as follows. `package-utility.sh` is a bash script which iterates each line of L1 and determines the behavior according to the file type (common files, directories, symbolic link files) and the

```
the path of package is: /tmp/package-hep

the total size of package is: 21GB

the total number of files is: 15881

the total number of directory is: 1548

the total number of symbolic link is: 4661
```

Figure 6: Package Description Infomation

system call type. The package description information is shown in Figure 6.

```
parrot_run /bin/bash package-utility.sh -L namelist
```

Each file inside L1 except for common dependencies is copied into the package, the final path of the file becomes the path of the package, followed by the original file path. In this program, the path of the package is **/tmp/package-hep**, so the final path of one file with the original path path1 is **/tmp/package-hep/path1**.

The map relationship between the file access path used in the actual analysis program and the actual file location used during the reproduction process is kept inside the map file, as illustrated in Table 5.

To ensure the successful reproduction, the filesystem structure of the original execution environment should be pre-

Path used in Program	Actual Location
/	/tmp/package-hep
/tmp/package-hep	/tmp/package-hep
/dev	/dev
/misc	/misc
/net	/net
/proc	/proc
/sys	/sys
/var	/var
/selinux	/selinux

Table 5: Structure of Map File

served as completely as possible. However, attempting to copy the whole content of one directory or one file is space-consuming and time-consuming, because the original program may only access the metadata of one file with the size of 200GB. Our solution is to determine the copy degree of one directory or one common file according to the type of the system call.

(3) rerun the analysis program using the package

When another scholar wants to repeat the analysis program, the description of the exepriment environment is firstly referred to build the underlying OS and basic software environment. In this example, the only necessary software basis is Parrot. Then the scholar can obtain one copy of the package and its map file, and rerun the program using the following command:

```
parrot_run -m /tmp/mountlist /bin/tcsh script_v4.sh
```

4.3 Structure of Map File
The map file of one package supports pattern matching semantics, which greatly reduces the scale of the map file and improves the efficiency of file path redirection. For example, one item inside the map file is /dir1 /package-dir1 means that the access to each file and subdirectory under /dir1 will be redirected into /package-dir1. If the semantics of the map file do not support pattern matching, the size of the map file will become astonishing. /dir1 may contain thousands of files, which will generate thousands of map items and take more time to redirect file paths.

During the packaging process, we notice that it is impossible to copy the files under certain directories. For example, during namelist acquisition process, three special files under /dev directory were recorded into the namelist: tty, null and urandom, which are used for input and output. Trying to copy these files into the new package would cause the packaging utility halt. We also notice that copying the files under directories like /proc is meaningless. Because the process id of one program is random and depends on the current allocation of process ID.

As for these special file paths, letting the program directly utilize the files on the new machine, where the new user repeats the program, is a better choice. The semantics is implemented by setting the target path of one file equal to the origin file path inside the map file. For example, /dev /dev means that when the program needs to access

files under /dev directory, it will directly use the file under /dev on the new machine, which is independent from the namespace of the package. This semantics also make the reference of files outside the package possible. If the new user wants to expand the analysis of one program to his own data with the path of /A, he can add the analysis code into the analysis script and add one item /A /A into the map file.

4.4 Discussion of Solution 4
Under Solution 4, the reproduction of one analysis program becomes easier. Rethink Matthias's example, the reproduction only needs to set necessary environment variables, the actual analysis command and the package. If different scholars want to repeat one analysis program, what they need to do is to obtain the package and rerun the actual analysis program. Under Solution 2, each scholar needs to get the necessary data and software, and then prepare software environment.

The starting point of Solution 4 is one successful execution on the original machine. Parrot traps all the system calls and copies each accessed file (except for common dependencies) into one package. The necessary scientific data and software dependencies will be copied into the target package. Strictly speaking, the concept of software and the concept of the scientific data are lost in Solution 4, because the granularity of parrot is file. The same thing happens to networked resources. Packaging utlity maintains one list of common dependencies and first judges whether the file belongs to this list before trying to copy one file into the package.

5. EVALUATION
The section first illustrates the relationship of different preservation and reproduction solutions, then the platform used to evaluate different solutions is introduced. The execution time and data size of Solution 2 and 4 are tested and compared. It is hard to accurately calculate the time consumption of Solution 1. Solution 3 shares the same time consumption and data size with Solution 3.

5.1 Relationship of Different Solutions
The relationship of these four solutions to repeat one program is shown in Table 6. Each new solution tries to make it easier to repeat one program through making the data and software preparation easier.

5.2 Evaluation Platform
The first three solutions are all directly verified on the machine the original author used to execute the experiment, which is convenient but involves access authority problem. Solution 4 is verified on the original machine and on one virtual machine. To compare the execution time and data size of Solution 4 with Solution 2, the package generated on the original machine is utilized to rerun the experiment on the original machine. This choice avoids the impact of different CPUs of the original machine and the new machine to repeat the experiment. To completely verify the correctness of Solution 4, one virtual machine sharing the same kernel version with the original machine is deployed. Then the necessary software environment is configured on the virtual machine,

Solution ID	Data Resources	Software Resources	Operation Manual
Solution 1	grid	CMSSW Git HTTP	email
Solution 2 and 3	local (HDFS)	CMSSW Git HTTP	shell script
Solution 4	package	package	user manual of package

Table 6: The relationship of different solutions

Sub-Task	Time	Percentage
Software acquisition from CMSSW	7min 24s	21.44%
Software acquisition from Git	9s	0.43%
Software acquisition from Wget	38s	1.83%
Environment Build - SCRAM	5min 48s	16.80%
Environment Build - Python	1s	0.05%
Data analysis	20min 31s	59.44%
Total	34min 31s	100.00%

Table 7: Breakdown of Execution Time of Solution 2

Data Category	Location	Data Size
HDFS	hadoop	20GB
PanFS	pscratch	1.6MB
CVMFS	cvmfs	103MB
NFS	opt	53KB
AFS	afs	34MB
Local FS	bin etc lib lib64 sbin tmp usr	68.3MB
Total		21GB

Table 10: Data Size of Solution 4

Data Category	Data Size
Software from CMSSW	448.3MB
Software from Git	73.7MB
Software from Wget	52MB
HDFS	43.3GB
Total	43.86GB

Table 8: Data Size of Solution 2

Data Category	Solution2: Size	Solution4: Size
HDFS	43.3GB	20GB
PanFS	155TB	1.6MB
CVMFS	7.4GB	103MB
NFS	862GB	53KB
AFS	10.2GB	34MB
Local FS	110GB	68.3MB
Total	156TB	21GB

Table 11: Data Size Comparison beweeen Solution 2 and 4

and the package together with the map file are copied onto it, and the experiment is repeated on the virtual machine.

5.3 Evaluation of Solution 2
The breakdown of execution time of Solution 2 is shown in Table 7. About 40% of the total execution time is consumed to prepare relevant software environment. The cvs command is called to check out 23 packages from CMSSW and three git repositories are cloned.

This solution does not support mining of implicit dependencies. The size of data and software from each category is shown in table 8, which only lists the size of each clear dependency.

5.4 Evaluation of Solution 4
The breakdown of execution time of the Solution 4 is illustrated in Table 9. The time used to obtain filename list and generate P1 is greatly longer than the execution time within the new package. However, the time consumption of filename list acquisition and package generation is one-time. That is, once the package is generated, many users can directly obtain the package and repeat the experiment separately.

Sub-Task	Time
Obtain file namelist	28min 23s
Generate P1	85min 51s
Re-run the program within P1	13min 4s

Table 9: Time Breakdown of Solution 4

The data size of Solution 4 is shown in Table 10. The data categories are more complex than our imagination. Except the data stored in Hadoop and the necessary set of software, files from local filesystem and several remote filesystems, such as PanFS, CVMFS, NFS, and AFS, are also necessary for the reproduction of the program.

5.5 Data size Comparison of Solution 2 and 4
The packaging utility checks the system call of each file within the namelist, and maintains the minimum dataset, which makes the size of the final package as small as possible. The total size of Hadoop under Solution 2 is astonishing, while the size of Hadoop under the package is decreased to 20GB. The total size of PanFS is 155TB, however, the size of the really used data of PanFS is only 1.6MB. Table 11 compares the original size and the really used size of each data source. Solution 4 reduces the total size of all data sources from 156TB to 21GB.

5.6 Execution Time Comparison of Solution 2 and 4
Table 12 shows the execution time comparison between Solution 2 and 4. The time consumption of the reproduction of one program from different scholars kept the same under Solution 2, including software and data preparation. However, under Solution 4, the data and software preparation is one-time. The following reproduction of the same program only needs to obtain one copy of the package and execute the actual experimental analysis directly. Data is obtained through accessing HDFS in Solution 2, but is copied into

Task Category	Solution 2: Time	Solution 4: Time
Software Acquisition	8min 11s	N/A
Environment Build	5min 49s	3s
Obtain file namelist	N/A	28min 23s
Generate package	N/A	85min 51s
Actual Analysis	20min 31s	13min 4s

**Table 12: Execution Time Comparison between Solution 2 and 4**

Figure 7: Preservation Integration of Multiple Programs

the package in Solution 4. This localization of experimental data speeds up the data analysis process, resulting the actual analysis time reducing from 20 minutes to 13 minutes.

6. DISCUSSION OF DATA AND SOFTWARE PRESERVATION

This section concludes the important lessons of data and software preservation and experiment reproduction learnt from this case study. The necessity of preservation integration of multiple programs, preservation granularity, how to preserve large data, how to mine data and software dependencies, the presrvation degree of the original filesystems, the importance of investigating user access model of data and software, will be discussed.

6.1 Preservation Integration of Multiple Programs

Different experimenters may execute different analysis programs on the same dataset using the same or overlapping sets of software. Generating one new package for each analysis program from scratch is time-consuming and space-consuming, which makes one data and software preservation mechanism, that can integrate the preservation requirements of different analysis programs, become necessary.

To integrate the data and software from multiple analysis programs into one archive, the concrete information of data and software, such as size, version, source and the number of files, need to be recorded. The packaging process needs to be re-organized. Packaging utility needs to maintain one list of current software and data subset already stored in the archive. When one user utilizes the packaging utility to preserve one program, the packaging utility scans the script, gets the data and software dependency part, judges whether each dependency has been inside the package, and only adds the new data and software into the archive, and updates the package information and relavant retrieval information.

The architecture of preservation integration of multiple programs is shown in Figure 7. Script 1 and Script 2, sharing one dependency B, belongs two different programs. Without preservation integration, B will be preserved twice in the archive. To improve the utilization efficiency of storage resources, preservation integration of multiple programs is necessary. In the archive, only one copy of B is stored, which is referred by both scripts.

6.2 Preservation Granularity

Another important factor of data and software preservation is the choice of preservation granularity. In our implementation of Solution 4, the preservation granularity is one file, because Parrot virtual filesystem traps each system call and redirects the access path of each file. As a result, during the packaging process, each time only one file can be copied into the target package, which is low-efficient. However, there are other options of preservation granularity. The software is generally preserved in the unit of package inside the remote repository, the packaging process of one experiment can adopt one package as the unit. In addition, if the size of one whole software repository is small and the access frequency of each package inside it is high, the granularity can be set to the whole repository.

6.3 Preservation of Large Data

The preservation policy must concern about the size of data. If data size is small, copying it into one pacakge and transferring it between different machines is not a bad idea. However, when the size becomes large, copying it into one package will make the size of the package uncontrollable. In Figure 7, the size of C is very large and easily to be obtained through Internet by different users. Directly referring C as networked resources is more efficient than preserving it into one package.

6.4 Mining of Dependencies

The mining degree of data and software dependencies greatly determines the repeatability of one experiment. Clear data dependencies is easy to figure out, but how to figure out implicit data dependencies is complex. The situation of software dependencies is more complex. The software dependencies include OS, hardware platform, direct software dependencies and indirect software dependencies. The first three categories are easy to deal with. The acquisition of indirect software dependencies needs more efforts.

One solution is to treat the computing environment, software environment and scientific data as one integral entirety and preserve the entirety completely. The reproduction of one experiment using this solution is easier. However, the original machine is not only for one special experiment and may concern a large amount of software and data unrelevant to the experiment. The time and space overhead of this solution must be considered.

Another solution is to utlize package management tools, like rpm, pacman, apt, to recursively trace each direct software dependency and generate one clear dependency graph. With the dependency graph, the new user can configure the software environment and repeat the experiment.

The third solution is to trace the system calls of the experiment execution process and record the paths of accessed files, which including the files of software dependencies. Parrot adopts this solution. Even if git is one software dependency, there is no necessary to install git on the new machine. All the files to guarantee the accurate execution of git has been copied into the package.

6.5 Preservation Degree of Directory Structure

Data and software preservation of physics experiments must concern the preservation degree of the original filesystem. The three core components of one filesytem are file, directory and metadata. The access of files can be divided into two categories: only access file metadata, like size, ownership, access authority list, and access both the file content and metadata. If both the content and metadata need to be accessed, the whole file must be copies into the package. However, if only the metadata of one file is accessed, how to preserve the file needs to be considered carefully. If the file size is 1KB or 1MB, preserving the content is acceptable. If the size of one file is 1GB, preserving the whole file just for accessing its metadata is inefficient.

Directory is one list of files and subdirectories under it. At first, we only copy the accessed items under one directory and ignored the unaccessed items. However, some operations on one directory depends on each item under it. For example, `ls /dirA` is one linux command to list the name of each item under `/dirA`. Because the directory itself has the name list of items under it, only the path of the directory is recorded into the namelist in the first phrase of Solution 4. However, when we repeat the command `ls /dirA` using the preserved package, it returns NULL because none of items under it exists in the package, the behavior of the whole program will become incorrect and unpredictable.

To ensure the semantics of the original program, all the items under one directory must be copied into the package. However, the time and space overhead appears again. Currently, Solution 4 creates one item with the same name but the size of zero for each item under `/dirA`. Another solution is to introduce one database to preserve the metadata of each file and directory.

6.6 User Access Model of Data and Software

The user access module of remote data and software has a great impact of the preservation mechanism, especially when we try to integrate the preservation of multiple programs. If all the programs only refer the original data and software and refuse to modify them, the preservation is easy and can ignore the data version problem. However, if some programs try to modify the original data and software to accustom to its own requirement, the preservation mechanism must provide one way to differentiate the original data and the special data used in one certain program. To have a clear understanding of the data access model, more examples need to be investigated.

6.7 Here is the Useful Package

At the beginnig of our case study, we generated one package for Matthias's example based on Solution 4. Several months later, the original machine met some problems and the access model of CMSSW was also modified. In this case, if you want to repeat Matthias's example on the original machine, different configurations, including CMSSW_ARCH, environment variables and CVS access authority need to be modified. However, with the generated package and its map file, we can repeat the experiment directly on the same machine without any modification.

7. RELATED WORK

Generally, there are three approaches to preserve software environment: hardware preservation, migration and emulation. Hardware preservation preserves the original software and its original operating environment. Software migration technique [7, 15] was used to facilitate running software on new machines. However, migration often involves the re-compiling and re-configuring the source code to accustom a new hardware platform and software environment. Emulation recreates the original software and hardware environment by programming future platforms and OSs. One common solution to implement this is virtual machine. According to the usage and emulation degree of the real machine, virtual machine can be divided into system virtual machine and process virtual machine. The working principle, design principle and performance evaluation of system virtual machine were illustrated in [10, 23]. The functionality of system VM to support different guest operating systems was illustrated in [2, 13, 21]. F. Esquembre [9] illustrated how JVM, one process virtual machine, can expedite the creation of scientific simulations in Java. The pros and cons of these three approaches were discussed in [18, 19, 11].

The preservation of computing environment and software environment was treated as one entirety in [18, 19, 11]. However, frequently changing experiment software makes the maintenance of the preserved experimental environment very complex. CernVM [5] treated them as two different categories. The preservation of computing environment is implemented with CernVM, and the preservation of software environment is based on a CernVM filesystem(CVMFS) specifically designed for efficient software distribution.

The importance of preserving software in source code format was emphasized in [25, 6]. However, CVMFS [5] published pre-built and configured experiment software releases to avoid repeating the time-consuming software building procedure.

Attempts from different perspectives to facilitate the reproduction of scientific experiments utilizing preserved software library has been made. The software distribution mechanism over network was discussed in [8, 3]. J. R. Rice et al. [20] made the reproduction process easier through the integration of user interface, scientific software libraries, knowledge base into problem-solving environment. S. R. Kohn et al. [14] tried to enable the creation and distribution of language-independent software library by addressing language interoperability. a scalable, distributed and dynamic workflow system for digitization processes was proposed in [22]. A distributed archival network was designed

in [24] to facilitate process-oriented automatic long-term digital preservation. M. Agosti et al. [1] aimed to help non-domain users to utilize the digital archive system developed for domain experts.

Current mechanisms of preserving scientific experiments assume that all the data and software mentioned in the experiments are necessary for the reproduction of the experiments. However, this is not always right. In some cases, the original author may leave additional code referring to irrelative data and software in the experiment programs. One mechanism, which can figure out the absolutely relevant data and software of one experiment, is important for both the preservation and reproduction of scientific experiments.

B. Matthews et al. [16] introduced one conceptual framework for software preservation from several case studies of software preservation. One tool to capture software preservation properties within a software environment was designed in [17] through a series of case studies conducted to evaluate the software preservation framework. L. R. Johnston et al. [12] proposed one overall data curation workflow for 3-5 case studies of preserving research data. Two case studies [4] were conducted to figure out the properties of data to be reused in the future, including type, purpose, new users. To figure out how to preserve HEP applications, this paper studies one case of preserving one representative HEP application.

8. CONCLUSION
In this paper, we tries to illustrate the challenges involved in data and software preservation and reproduction through one case study in preserving one High Energy Physics application. Four different solutions to repeat the original application, together with the pros and cons of each one, are proposed and analyzed. Each new solution aims to improve the reproductivity of the application and reduce the complexity of reproduction. Finally, we propose one fine-grained packaging toolkit to help the original author to generate one independent package containing all the necessary dependencies except for common dependencies, and repeat the experiment from scratch on two Amazon EC2 virtual machines successfully.

In the future, more case studies will be investigated to verify the correctness of our packaging toolkit. The challenges we propose in this paper will also be investigated.

9. REFERENCES

[1] M. Agosti and N. Orio. To envisage and design the transition from a digital archive system developed for domain experts to one for non-domain users. In *Proceedings of the 12th ACM/IEEE-CS joint conference on Digital Libraries*, pages 11–14. ACM, 2012.

[2] P. Barham, B. Dragovic, K. Fraser, S. Hand, T. Harris, A. Ho, R. Neugebauer, I. Pratt, and A. Warfield. Xen and the art of virtualization. *ACM SIGOPS Operating Systems Review*, 37(5):164–177, 2003.

[3] J. Blomer, P. Buncic, and T. Fuhrmann. Cernvm-fs: delivering scientific software to globally distributed computing resources. In *Proceedings of the first international workshop on Network-aware data management*, pages 49–56. ACM, 2011.

[4] C. L. Borgman. Data, data use, and scientific inquiry: Two case studies of data practices. In *Proceedings of the 12th ACM/IEEE-CS joint conference on Digital Libraries*, pages 19–22, 2012.

[5] P. Buncic, C. A. Sanchez, J. Blomer, L. Franco, A. Harutyunian, P. Mato, and Y. Yao. Cernvm–a virtual software appliance for lhc applications. In *Journal of Physics: Conference Series*, volume 219, page 042003. IOP Publishing, 2010.

[6] M. Castagné. Consider the source: The value of source code to digital preservation strategies. *SLIS Student Research Journal*, 2(2):5, 2013.

[7] C. Cifuentes and V. Malhotra. Binary translation: Static, dynamic, retargetable? In *Software Maintenance 1996, Proceedings., International Conference on*, pages 340–349. IEEE, 1996.

[8] G. Compostella, S. P. Griso, D. Lucchesi, I. Sfiligoi, and D. Thain. Cdf software distribution on the grid using parrot. In *Journal of Physics: Conference Series*, volume 219, page 062009. IOP Publishing, 2010.

[9] F. Esquembre. Easy java simulations: A software tool to create scientific simulations in java. *Computer Physics Communications*, 156(2):199–204, 2004.

[10] R. P. Goldberg. Survey of virtual machine research. *Computer*, 7(6):34–45, 1974.

[11] N. C. Hong, S. Crouch, S. Hettrick, T. Parkinson, and M. Shreeve. Software preservation benefits framework. *Software Sustainability Institute Technical Report*, 2010.

[12] L. R. Johnston. A workflow model for curating research data in the university of minnesota libraries: Report from the 2013 data curation pilot. 2014.

[13] A. Kivity, Y. Kamay, D. Laor, U. Lublin, and A. Liguori. kvm: the linux virtual machine monitor. In *Proceedings of the Linux Symposium*, volume 1, pages 225–230, 2007.

[14] S. R. Kohn, G. Kumfert, J. F. Painter, and C. J. Ribbens. Divorcing language dependencies from a scientific software library. In *PPSC*, 2001.

[15] D. Mancl. Refactoring for software migration. *Communications Magazine, IEEE*, 39(10):88–93, 2001.

[16] B. Matthews, B. McIlwrath, D. Giaretta, and E. Conway. The significant properties of software: A study. *JISC report, March*, 2008.

[17] B. Matthews, A. Shaon, J. Bicarregui, and C. Jones. A framework for software preservation. *International Journal of Digital Curation*, 5(1):91–105, 2010.

[18] B. Matthews, A. Shaon, J. Bicarregui, C. Jones, J. Woodcock, and E. Conway. Towards a methodology for software preservation. 2009.

[19] T. A. Phelps and P. B. Watry. A no-compromises architecture for digital document preservation. In *Research and Advanced Technology for Digital Libraries*, pages 266–277. Springer, 2005.

[20] J. R. Rice and R. F. Boisvert. From scientific software libraries to problem-solving environments. *Computational Science & Engineering, IEEE*, 3(3):44–53, 1996.

[21] M. Rosenblum. VmwareâĂŹs virtual platformâĎć. In

Proceedings of Hot Chips, pages 185–196, 1999.

[22] H. Schöneberg, H.-G. Schmidt, and W. Höhn. A scalable, distributed and dynamic workflow system for digitization processes. In *Proceedings of the 13th ACM/IEEE-CS joint conference on Digital libraries*, pages 359–362. ACM, 2013.

[23] J. E. Smith and R. Nair. The architecture of virtual machines. *Computer*, 38(5):32–38, 2005.

[24] I. Subotic, L. Rosenthaler, and H. Schuldt. A distributed archival network for process-oriented autonomic long-term digital preservation. In *Proceedings of the 13th ACM/IEEE-CS joint conference on Digital libraries*, pages 29–38. ACM, 2013.

[25] J. G. Zabolitzky. Preserving software: Why and how. *Iterations: An Interdisciplinary Journal of Software History*, 1(13):1–8, 2002.